

AI CONSULTING · FINAL PROJECT

# Taste Agent for Letterboxd

From viewing history to semantic taste  
intelligence

— Marc



# Letterboxd already owns an unusually rich taste dataset.

## MEMBERS




## EXISTING BEHAVIORAL ASSETS

- Ratings
- Viewing history
- Watchlists
- Reviews
- Likes

The opportunity is not collecting more taste data. It is **extracting more intelligence** from data Letterboxd already has.

Paid memberships are a core part of Letterboxd's existing business model.

# Traditional metadata explains what we watch — not what characteristics drive preference.



Discovery can move from content similarity toward **preference similarity**.

Semantic signals identify characteristics associated with preference. They do not infer personality or identity.

# Turn behavioral history into a reusable semantic Taste Model.

## SIGNALS IN

- Long-term taste
- Recent taste
- Semantic preferences
- Cross-media patterns
- Current intent



## TASTE MODEL

The core AI asset —  
reusable across  
surfaces

## SURFACES IT COULD POWER

### My Taste

Visual explanation of preferences

### Semantic Discovery

Recommendations from preference  
patterns

### Taste Agent

Conversational film discovery

### Taste Evolution

How preferences change over time

Letterboxd · core signal


Goodreads · tested enrichment

Spotify · future opportunity

| The chatbot is not the product. The Taste Model is the core AI asset.

Letterboxd knows your taste in film. Can broader cultural taste make it better at discovering your next film?

# I built the semantic layer as a working no/low-code POC.



**107**  
VALID WORKS  
96 films + 11 books

**62**  
SEMANTIC  
DIMENSIONS

**2**  
MEDIA TYPES

The LLM classifies the **work** — not the **user**.

Ratings and preference signals are kept outside the LLM classification step.

BUILT WITH

n8n

TMDB



Google Books

OpenAI

JavaScript

Python migration planned for deterministic processing and recommendation evaluation in the MVP.

# The POC evolved through experimentation — not one successful prompt.



Seeing a characteristic frequently does not necessarily mean it predicts **stronger preference**.


## PROVEN

Semantic preference modelling is technically feasible.

## NOT YET PROVEN

That it produces better recommendations.

# From synthetic viewing history to a structured Taste Model.



## EXAMPLE TASTE MODEL OUTPUT — TOP PREFERENCE ASSOCIATIONS



The POC demonstrates semantic preference modelling and cross-media technical compatibility — not recommendation quality.

Example values are from the development validation run and demonstrate the structure of the Taste Model.

# Taste intelligence only matters if it creates measurable product value.

FREE → PRO

## Convert

A stronger personalized discovery experience creates an additional reason to upgrade.

PAID → RENEW

## Retain

Better discovery can increase recurring value from a paid membership.

PRO → PATRON

## Upsell

Advanced Taste Evolution and optional cross-media capabilities could support higher-tier differentiation.

### HYPOTHETICAL PACKAGING TO TEST

#### Free

Basic Taste preview

#### Pro

Full Taste Model  
Semantic recommendations  
Taste Agent

#### Patron

Taste Evolution  
External cultural connections  
Cross-media model

| Monetize **within the existing membership model** rather than create a separate AI subscription.

Illustrative product hypothesis — not current Letterboxd subscription packaging.

# The commercial hurdle is measurable — and testable.

≈ 3,948

Incremental \$19 Pro-subscriptions required to recover an illustrative \$75K Year-1 investment.

≈ 0.08 percentage points

of incremental conversion across an illustrative 5M relevant Free-member cohort.

## ILLUSTRATIVE YEAR-1 SCENARIOS

\$101K

Conservative  
35% ROI

\$218K

Base  
190% ROI

\$483K

Upside  
543% ROI

**SCENARIO ANALYSIS — NOT A LETTERBOXD REVENUE FORECAST**

The Pilot should test whether product value crosses the break-even threshold — not assume that it will.

Actual subscriber segmentation, conversion and churn are not publicly disclosed.

# The biggest risks are product and commercial — not whether the LLM can produce JSON.

20/25

## RECOMMENDATION VALUE

Semantic recommendations may not outperform the baseline.

20/25

## COMMERCIAL VALUE

Better recommendations may not change subscription behavior.

15/25

## GDPR / PROFILING

Consumption history and the inferred Taste Model are personal data.

12/25

## METADATA QUALITY

Incorrect matches can contaminate semantic vectors.

### EU AI ACT

**Entertainment recommendation:** not high-risk under the proposed intended use.

**Conversational AI:** transparency requirements apply.

### GDPR

Purpose limitation

User transparency & control

Data minimization


Optional cross-media connections

DPIA before production

Understand what characteristics a user **responds to in culture** — not who that user is as a person.

## RECOMMENDATION

**The POC passes the technical gate. The next investment should test recommendation value.**



**Don't fund full deployment yet. Fund the experiment that determines whether Taste Agent discovers a better next film.**

Technical feasibility → Product evidence → Commercial evidence → Scale